



Andres Torrubia

✉ antorbus@gmail.com

🌐 andrestorrubia.com

🐙 github.com/antorbus

☎ +1 919 918 0912

🌐 [linkedin](https://www.linkedin.com/in/antorbus)

Even if Moore's Law dies, I believe there are still orders of magnitude of performance waiting to be extracted from ingenuous fabrication of the whole tech stack, from silicon to API.

EDUCATION

Duke University

Triple Major in

- BSE Electrical and Computer Engineering
- BS Mathematics
- BS Computer Science

2023 - 2027

GPA: 3.89

SKILLS

Programming Languages

C/C++, CUDA, Python, Verilog, MIPS/RISCV assembly.

Languages

English (Native), Spanish (Native), Mandarin (C1).

Software

Proficient: PyTorch, NumPy, OpenMP, Git, Linux/Unix, Arduino.

Working Knowledge: KiCad, LTspice, Docker, Xilinx Vivado, GNU/LLVM toolchain.

Practical

Prototyping, soldering, data-sheet and ISA reading.

EXPERIENCE

HummingBird: A Fast IO Aware Attention Kernel For Small Sequences

Research (advised by Prof. Emily Wenger)

Fall 2024

- Built a multi-head self-attention kernel for small sequences (≤ 48 tokens).
- Achieved 3x wall-clock time speedup compared to state-of-the-art methods such as flash attention.
- Optimized specifically for 4090s and 3090s using bare CUDA at FP16.

Metamaterial Inverse Design Problem

Research (with Prof. David R. Smith)

Ongoing (Started April)

- Currently developing GPU-accelerated Discrete Dipole Approximation algorithms to simulate and control electromagnetic responses for 3d metamaterials.

Kaggle Silver Medal, GoogleAI4Code

Deep Learning Competition

Summer 2022

- Competed solitary against 1135 teams, ultimately ranking among Kaggle Grandmasters.
- Developed and trained LLMs to automatically reconstruct the order of markdown and code cells in Python notebooks.
- Approximated the non-differentiable Kendall Tau loss function using softmax to enable gradient-based optimization, later refined the approach to better approximate Spearman's rho.
- Learned Rich Sutton's "The Bitter Lesson."

FPGA Raycaster Engine

Software/Hardware Project

Fall 2024

- Designed and built a custom 5-stage pipelined, RISC-esque CPU architecture implemented on a FPGA and developed a full raycasting engine in bare-metal assembly, optimized the memory architecture to achieve real-time performance.

Recommendation Systems and Image Generation

Intern, Freepik Company

Summer 2022 and Summer 2023

- Developed a transformer-based recommendation system trained autoregressively, used Faiss for similarity search and multi-GPU for fast training and inference.
- Engineered an image dataset generation pipeline from public domain media using FFmpeg and captioning models. Validated the dataset's efficacy by fine-tuning diffusion models with LoRA.

LightLemur, A Lightweight Tensor Differentiation Library

Software Project

Ongoing (Started January)

- Leading a team of 3 people to build a reverse mode tensor differentiation library.
- Achieving close to PyTorch performance on CPU using OpenMP and SIMD optimizations.

AM Radio

Hardware Project

Spring 25

- In my quest to learn electronics, I have built an analog AM radio from scratch, from circuit to PCB, without internet help, only data-sheets and fundamentals.